

# Introduction to Programming (CS 101)

Spring 2024



## Lecture 17:

Pointers again (dynamic allocation, arrays of pointers)

**Instructor:** Preethi Jyothi

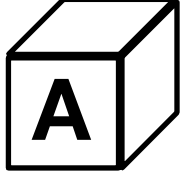
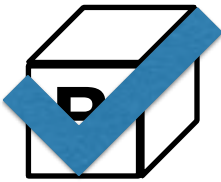
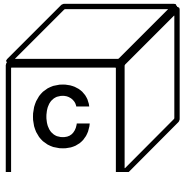
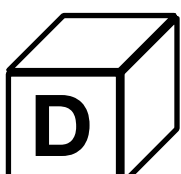
Based on material developed by Prof. Abhiram Ranade and Prof. Manoj Prabhakaran

# Recap (IA): Pointers vs. References

What is the output of the last `cout` statement?

```
#include <iostream>
using namespace std;

int main() {
    int i = 1, j = 2;
    int* p = &i, * q = &j;
    int& a = i, &b = j;
    a = b; i++; j++;
    cout << i << " " << j << endl;
}
```

	2	3
	3	3
	1	3
	3	2

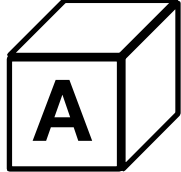
References **cannot be** reinitialized. `a = b` would behave the same as `i = j`

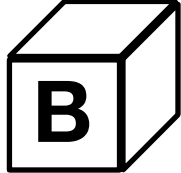
# Recap (IB): Pointers vs. References


What is the output of the last `cout` statement?

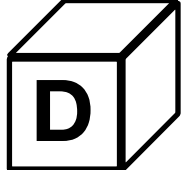
```
#include <iostream>
using namespace std;

int main() {
    int i = 1, j = 2;
    int* p = &i, *q = &j;
    int& a = i, &b = j;
    a = b; i++; j++;
    cout << i << " " << j << endl;
    p = q; (*p)++; (*q)++;
    cout << i << " " << j << endl;
}
```

 A 3 3

 B 4 4

 C 3 5

 D 2 3


 Pointers **can be** reinitialized. `p = q` would assign the address of `j` (`&j`) to `p` (replacing `&i`)

# Recap (II): Peril of Pointer Reassignment

What is the output of the following program?

```
#include <iostream>
using namespace std;
```

```
1  int main() {
2      int j = 1, i = 3;
3      int* a = &j;
4      int* b = a;
5      a = &i;
6      cout << *a << " " << *b << endl;
7  }
```



Say you have two pointers pointing to the same variable (line 4). Reassigning one of these pointers (a in line 5) will not reflect in the other (b).

# Recap (IIA): Pointer Arithmetic

What do the `cout` statements produce as output?

```
#include <iostream>
using namespace std;
```

```
int main() {
    int A[] = {1,2,3,4};
    int* p = A;
    p++; cout << *p << endl;
    cout << *(p++) << endl;
    cout << *(++p) << endl;
}
```

OUTPUT: 2

OUTPUT: 2

OUTPUT: 4



`*(p++)` and `*(++p)` deference from different addresses. `p++` (or `++p`) is done first, followed by dereferencing.

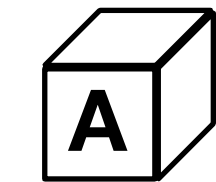
## Recap (IIB): Pointer Arithmetic

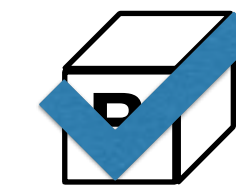
What is the output of this program?

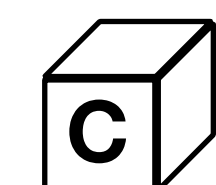
```
#include <iostream>
using namespace std;

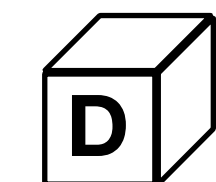
void sumslice(int* p, int* q) {
    int sum = 0;
    while(p < q) { sum += *p; p++; }
    cout << sum;
}

int main() {
    int A[] = {1,2,3,4};
    sumslice(A+1, A+4);
}
```

 5

 9

 UB (undefined behaviour)

 10



$p < q$  would evaluate to true if  $p + i < q + j$  where  $i$  is strictly less than  $j$

## Recap (IIC): Pointer Arithmetic

What is the output of this program?

```
#include <iostream>
using namespace std;

int main() {
    int A[] = {0,1,2,3};
    int* p = A;
    char* c = (char*)p;
    for(int i = 0; i < 4; i++)
        cout << int(*(c+i)) << " ";
    cout << endl;
}
```

OUTPUT: 0 0 0 0



Can cast an `int` pointer `p` to a `char` pointer `c`. Then, `c+i` would increment the pointer by 1 byte (i.e. size of a `char`). This casting allows us to print the integer byte-by-byte.



# Pointers in `struct` (continued)


## CS 101, 2025



# Linked list

- A linked list is a data structure consisting of nodes, where each node contains one or more data fields and a pointer to the next node in the list.

```
struct node {  
    int val;  
    node* next; // nullptr, if end of list  
};
```



- The first node of a linked list is typically referred to as the **head** node
- The last node of a linked list is the **tail** node, and its next node is a nullptr (indicating that the list has terminated)


# Reversing a linked list

Given the head of a list, reverse the list by changing the links between the nodes (without creating a new list)

```
node* reverseList(node* head) {  
    node* prev = nullptr;  
    node* curr = head;  
    node* next = nullptr;
```

```
    struct node {  
        int val;  
        node* next;  
    };
```

```
    while(curr != nullptr) {  
        next = curr->next;  
        curr->next = prev;  
        prev = curr;  
        curr = next;  
    }  
    return prev;  
};
```





# Dynamic allocation of memory (using pointers)

## CS 101, 2025

# Dynamically Allocated Memory

- Suppose we want to create a queue that can grow without limit (other than the limits set by the system policies/resources)
  - Create a queue that is as big as the maximum allowed?
  - But what if we want multiple such a priori unbounded queues?
- Ideally, the memory used for the queue should grow/shrink as the queue grows/shrinks
- More generally, we would like to create "boxes" in memory dynamically (decided at the time of program execution)


# A Bit about Memory

- Each program gets its own memory space
  - Memory isolation
- Not all of the addressable space will be used by a process
  - Physical memory will not be allocated until the process needs it
  - Virtual memory
- Mapping virtual memory to physical memory is quite complex and is handled by the operating system and the hardware
  - The program only works with the virtual memory

64 bit address space has  
16 Exabytes (16 billion GB)

# A Bit about Memory

- The virtual memory space is divided into different segments to hold various things needed by the program
- Dynamically allocated memory comes from one such segment called the heap which can grow/shrink as needed



# A Bit about Memory

- The virtual memory space is divided into different segments to hold various things needed by the program
- Dynamically allocated memory comes from one such segment called the heap which can grow/shrink as needed



A typical layout of virtual memory



## A Box in the Heap

- A `new` expression can be used to create "boxes" in the heap memory dynamically (i.e., decided at the time of program execution)
- But how will we access this box without a variable name?
- `new` returns a *pointer to the box*
  - It is the programmer's responsibility to save/use that pointer appropriately (and not lose it)

```
int* p = new int; // creates a new int "box" without a variable name!
*p = 7;           // we can access the new box only through its address
p = new int;      // oops! the previous box has become inaccessible now!
```

Memory Leak!



# new and delete operators

- Whenever `new` is used to dynamically allocate a region of memory in the heap, you must use `delete` to deallocate the storage (i.e. return the storage to the heap)
- `new` and `delete` should be used together

```
int* p = new int; // creates a new int "box" without a variable name
      :
delete p;          // p's use is over. release the memory used for it!
```

- `delete` does not actually delete the pointer variable. It returns the allocated heap memory back to the operating system.
- After deleting, the pointer variable can be assigned a new value. E.g., `p = nullptr;`

# Variable Length Arrays

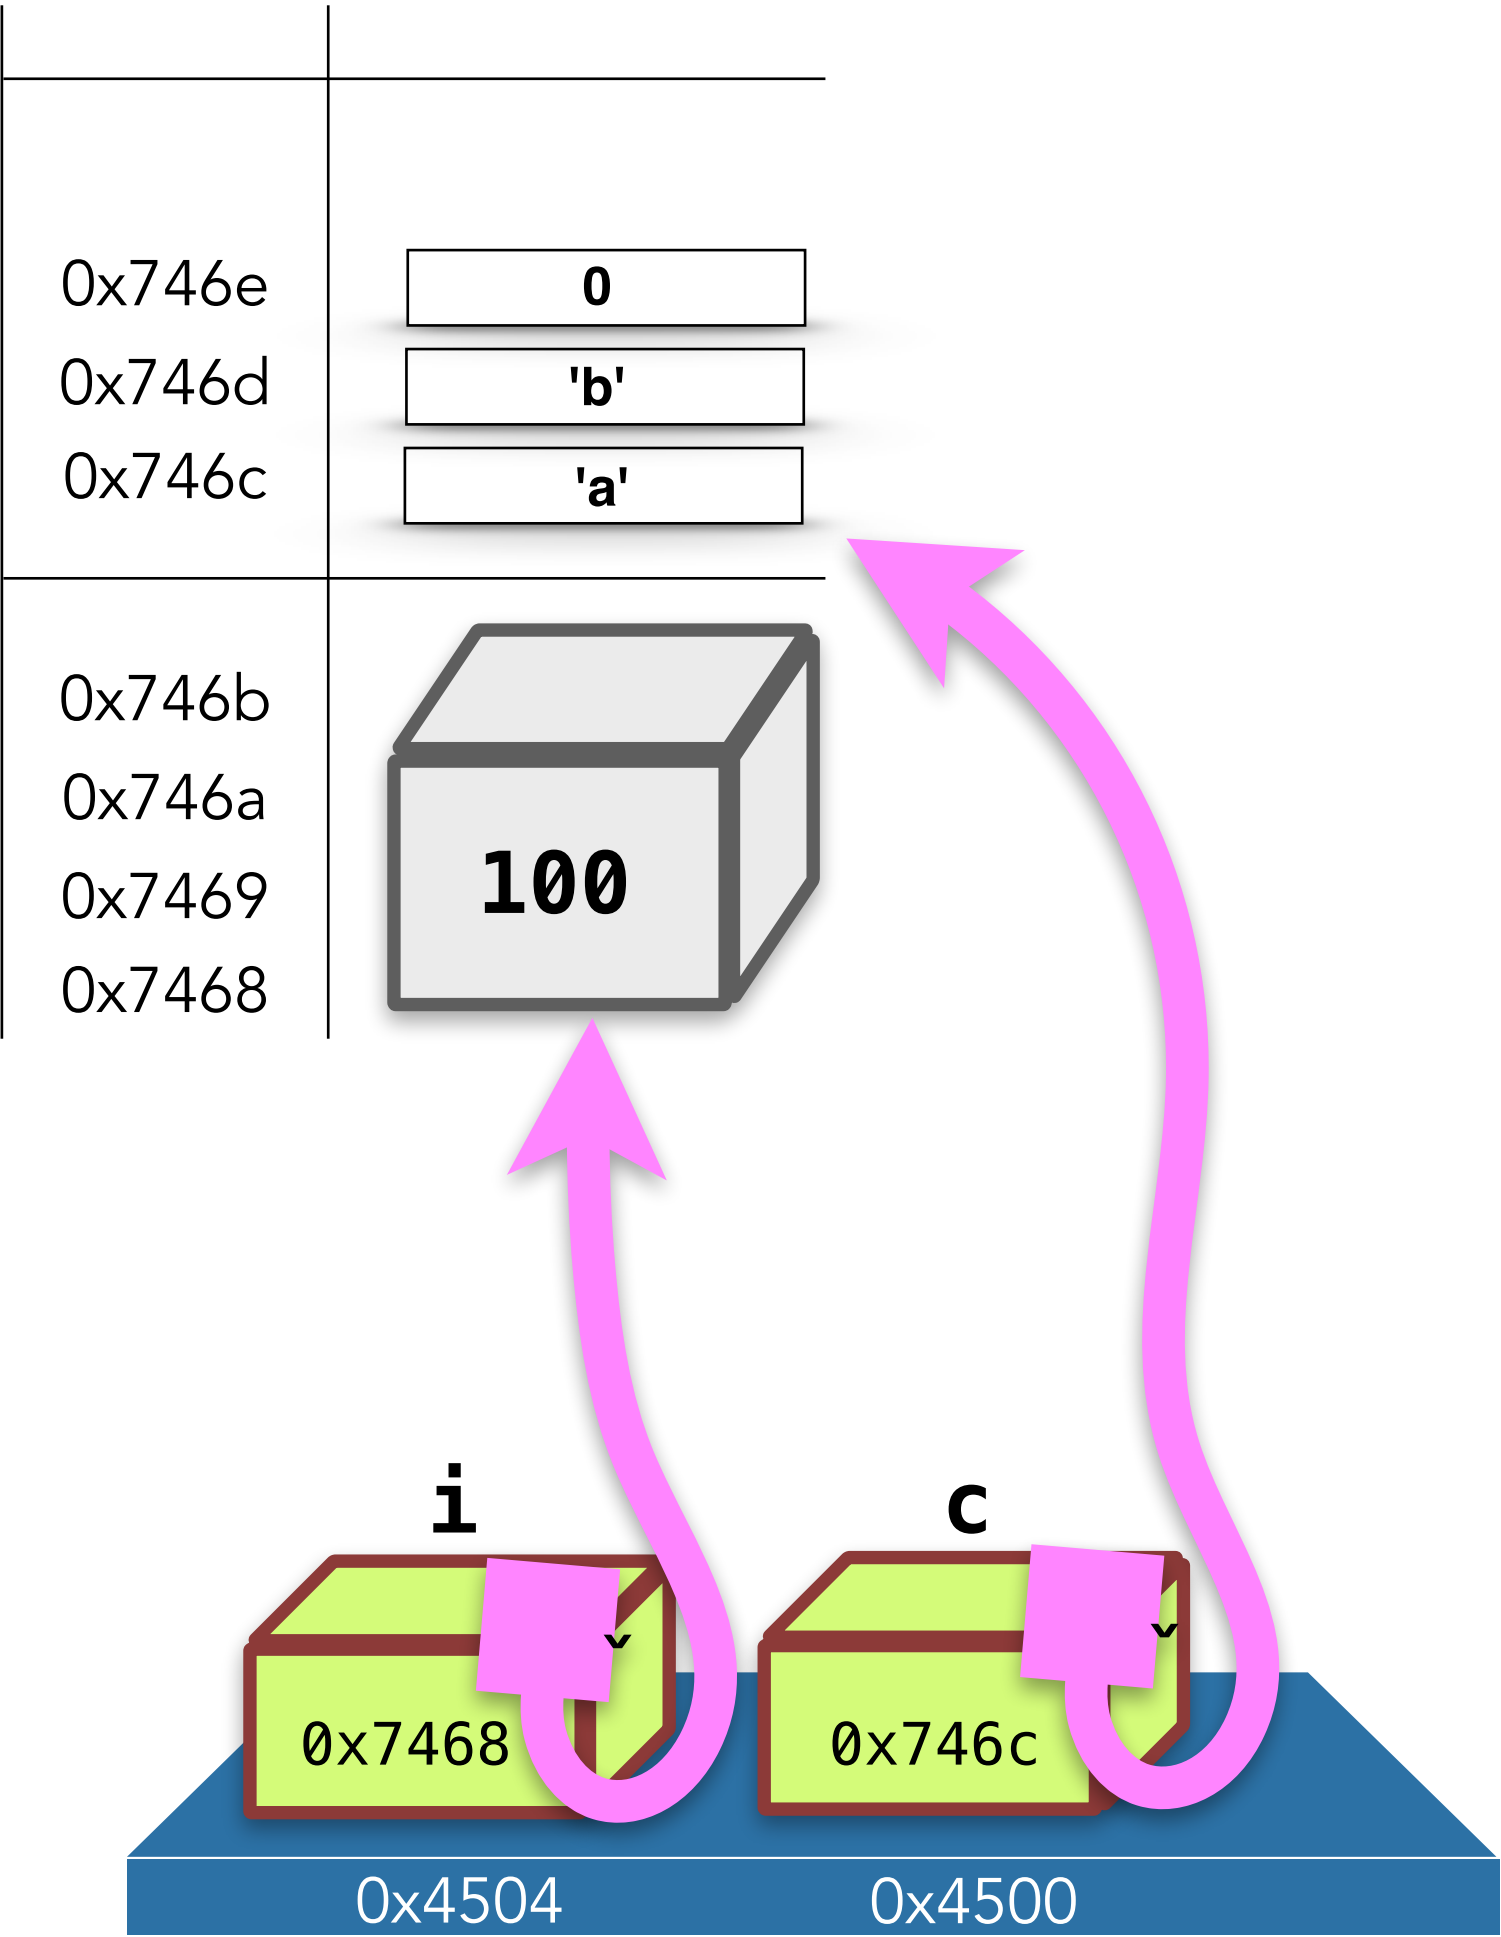
- Can create variable length arrays in the heap
- `new type[n]` expression returns a pointer to the first element in an array of  $n$  elements of the requested type
- Should free it using the operator `delete[]` (not using `delete` which will result in undefined behaviour)

```
unsigned n; cin >> n;
int* p = new int[n]; // p[0], ..., p[n-1] are allocated now
...                // use the array
delete[] p;         // release the memory for the entire array
p = nullptr;        // now we can overwrite the address
```

# Illustration of heap memory allocation using new and delete

```
int main() {
    int* i = new int;
    char* c = new char[3];
    *i = 100;
    c[0] = 'a';
    c[1] = 'b';
    c[2] = '\0';
    delete i;
    delete[] c;
}
```

Heap memory




Stack memory


# Pointer Bugs

Point out any bugs in this function. Need not be compiler errors.

p1, p3, p4 are not deallocated (deleted) 

```
void f() {  
    int* p1, *p2, *p3, *p4;  
    p1 = new int;  
    p3 = new int;  
    p4 = new int;  
    p3 = p1;  
    *p2 = 5;  
}
```

Memory leak! p3 = p1, means that we are overwriting p3 that used to contain the address of a variable in the heap. The latter can now no longer be addressed 😞 

p4 is allocated but never used. 

Dereferencing an uninitialized pointer 

## Example: Reading Inputs of Given Length

```
int* sort(int p[], int n); // returns an array allocated on the heap

int main() {
    int n; cin >> n;
    int* p = new int[n];
    for(int i=0; i<n; i++) cin >> p[i];
    int* q = sort(p,n);
    for(int i=0; i<n; i++) cout << q[i] << " ";
    cout << endl;
    delete[] p;           // release the memory allocated
    delete[] q;           // was allocated within sort as an array
}
```

## Example: Reading Inputs of Given Length

```
void sort(int in[], int out[], int n); // cleans up all new memory

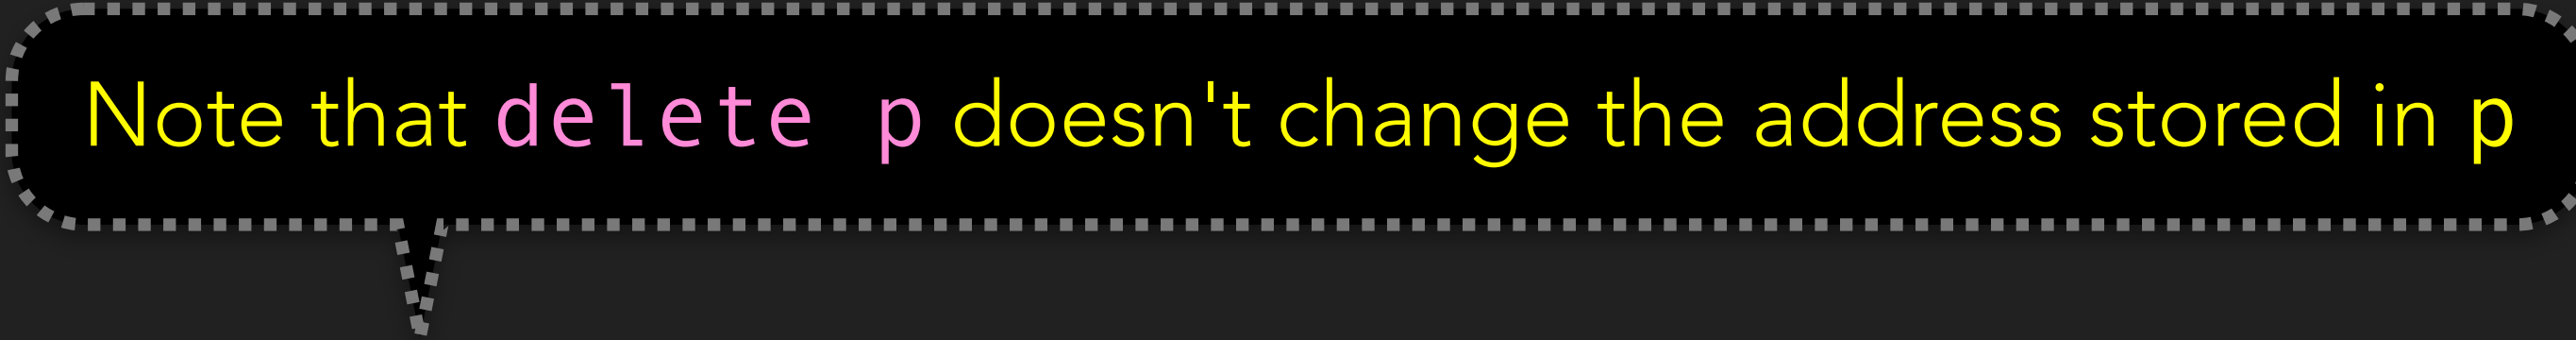
int main() {
    int n; cin >> n;
    int* p = new int[n]; int* q = new int[n]; // two arrays created here
    for(int i=0; i<n; i++) cin >> p[i];
    sort(p,q,n);
    for(int i=0; i<n; i++) cout << q[i] << " ";
    cout << endl;
    delete[] p;           // and two arrays deleted here
    delete[] q;           // easier to prevent memory leaks
}
```

## Some Tips

- Whenever you use `new` or `new[]` in your program, make sure there is a matching `delete` or `delete[]`
  - Even if it may look like it doesn't matter (small program, will anyway exit right after this,...), before exiting, a good C++ program should `delete` all the heap memory allocated via `new`
- Because `new` may do more than allocate memory and `delete` may do more than free it. (Maybe files or sockets are involved via a `new` invocation



## Some Tips

- Whenever you use `new` or `new[]` in your program, make sure there is a matching `delete` or `delete[]`  

- Accessing a deleted pointer is an error (undefined behaviour)
- Deleting an already deleted pointer is an error (crashes, typically)
  - Beware when multiple pointers may hold the same address
- C++ has several mechanisms to help with correctly using memory
  - Pre-implemented data structures in the standard library (e.g., vectors)
  - Constructor and destructor functions (coming up in a later class)





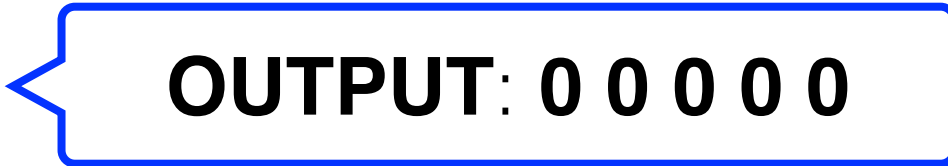
# Arrays of Pointers

## CS 101, 2025

# Arrays of pointers

- We can define an array of pointers, that is, each element of an array will be a pointer.  
Example:

```
int main() {  
    int* A[5];  
    int B[5] = {};  
    for(int i = 0; i < 5; i++) {  
        A[i] = &B[i];  
        cout << *A[i] << " ";  
    }  
}
```



# Arrays of pointers: Command line arguments to `main`

- So far, you run a C++ program on the command line using `./a.out`
- You can pass *additional arguments* to the `main` function using additional text after `./a.out`. E.g.:  

```
./a.out hello world
```
- This requires using an alternative overloaded definition of `main`, namely:  

```
int main(int argc, char* argv[]);
```
- Here, `main` takes two arguments:
  1. `argc`: `int` argument giving the number of space-delimited words typed on the command line
  2. `argv`: Array of pointers to C-style char arrays, with `argv[i]` containing the address of the  $i^{\text{th}}$  word on the command line

# Print command line arguments

```
int main(int argc, char* argv[]) {  
    for(int i = 0; i < argc; i++)  
        cout << argv[i] << endl;  
}
```

- If you use the following command: `./a.out hello world`

then, `argc = 3` (`./a.out`, `hello`, `world`) and `argv` would have three elements of type `char*` (pointing to null-terminated C-style `char` arrays) and it would print as output:

```
./a.out  
hello  
world
```